

StressPP: A Synthetic Workload Generator based on “Stress”

Politecnico di Milano, Advanced Computer Architectures Course

Andrea Guglielmini,
Matr. 893795



Why do we need a workload generator?

A workload generator is an useful tool to test the behaviour of an application while simulating different real-world hardware usages;

We are used to test programs while they are running alone on our computer but this environment is much different from the production one in which a number of other processes, with real workloads, may be running together with our application contending for resources.

This is why it's useful to test a target program in order to reason about its behaviour and its performance while integrated with possible other noisy neighbour processes.

StressPP





StressPP

StressPP is a synthetic workload generator based on [Stress](#).

It's main advantages are the usage of *threads* instead of processes and the capability to precisely stress a CPU core.

It's a flexible and ready to be applicable to real world scenarios.

It can be simply started via *command-line* or by writing your own script using it's straightforward APIs.



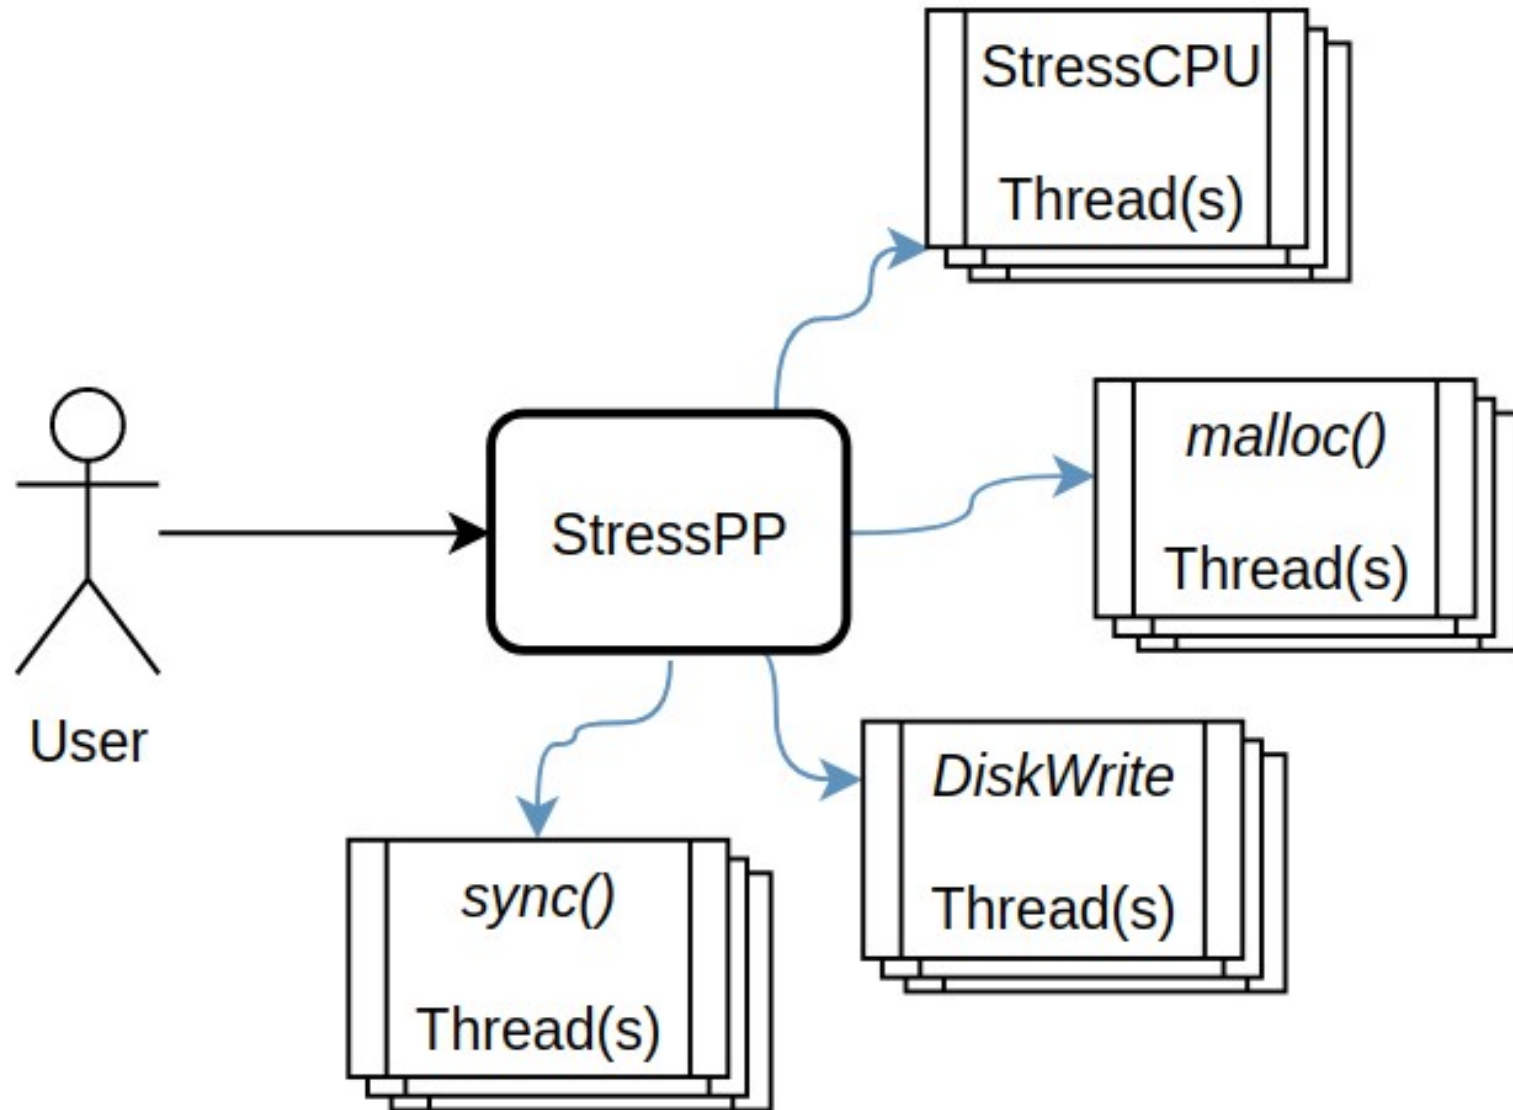
StressPP

It can compose three different types of workloads:


- **CPU Intensive:** issuing floating-point operations to designated CPU Cores.
- **Memory Intensive:** By allocating and random accessing virtual memory chunks.
- **Disk Intensive:** By writing and deleting fictitious files on local drive.
- **I/O Intensive:** By issuing several `sync()` to the underlying hardware.

The three modules can be executed independently or in chorus.

StressPP: Architecture



StressPP: Architecture




StressPP works by running different threads:

StressCPUThread is a wrapper for a plain *pthread*, it's used in order to easily *start, stop and manage the CPU affinity* of a single system thread.

An instance of *StressCpuThread* is instantiated for any core which the user requires to *stress* with floating point operations.

Optionally, the user may specify a CPU affinity matrix.

StressPP: Architecture



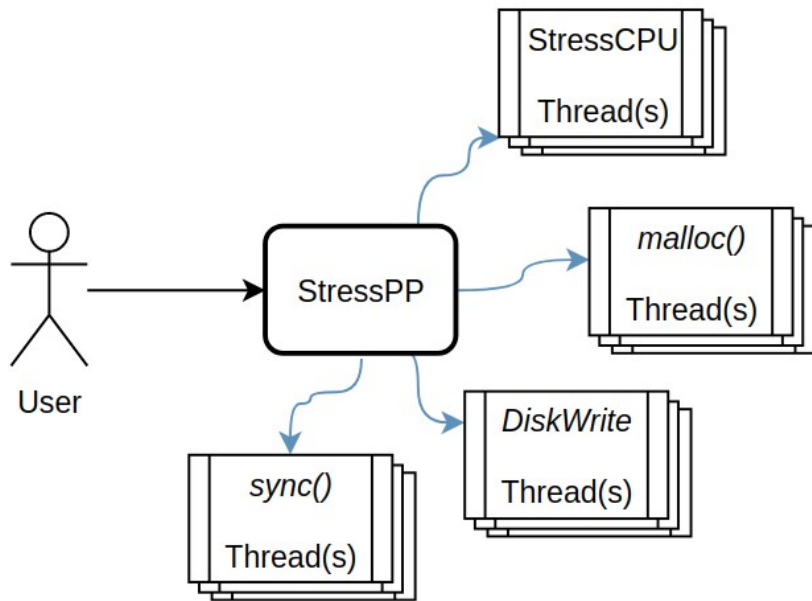
StressPP works by running different threads:

Malloc Threads are used in order to stress the virtual memory of the system.

The user may specify:

- How many threads to launch.
- How many *chunks* try to allocate.
- The size of a single chunk.

StressPP: Architecture




StressPP works by running different threads:

Disk Write Threads are used in order to stress the physical drive of the system.

The user may specify:

- How many threads to launch.
- How many *files* try to allocate.
- The size of a single file.
- If the allocated files needs to be deleted after completion.

StressPP: Architecture



StressPP works by running different threads:

Sync Threads are used in order to stress the I/O subsystem.

The user may just specify how many threads to launch.

StressPP: Validation

Metrics gathered while system in idle

avg-cpu:	%user	%nice	%system	%iowait	%steal	%idle		
	0,45	0,00	0,35	0,00	0,00	99,20		
Device	tps	kB_read/s	kB_wrtn/s	kB_dscd/s	kB_read	kB_wrtn	kB_dscd	
nvme0n1	1,00	0,00	20,00	0,00	0	100	0	
dm-0	5,20	0,00	20,80	0,00	0	104	0	
dm-1	0,00	0,00	0,00	0,00	0	0	0	

```
1  [| 0.7%]
2  [ 0.0%]
3  [| 1.3%]
4  [| 2.0%]
Mem[||||| 2.65G/15.4G]
Swp[ 0K/7.72G]
```

Performance counter stats for 'system wide':

2.033.097.949	cycles	#	0,017 GHz
55.126.405	cache-misses		
820.373.624	instructions	#	0,40 insn per cycle
120.313,74 msec	cpu-clock	#	4,000 CPUs utilized
22.942	page-faults	#	0,191 K/sec
148.889	r01C7	#	0,001 M/sec

30,078463556 seconds time elapsed

Note: r01C7 is the number of scalar double precision floating point operations

100

Metrics gathered while system in idle

avg-cpu: %user %nice %system %iowait %steal %idle							
0,45 0,00 0,35 0,00 0,00 99,20							
Device	tps	kB_read/s	kB_wrtn/s	kB_dscd/s	kB_read	kB_wrtn	kB_dscd
nvme0n1	1,00	0,00	20,00	0,00	0	100	0
dm-0	5,20	0,00	20,80	0,00	0	104	0
dm-1	0,00	0,00	0,00	0,00	0	0	0

Note: No program is actively writing on HDD.

```

1  [ | 0.7%
2  [ | 0.0%
3  [ | | 1.3%
4  [ | | 2.0%
Mem[||||| 2.65G/15.4G
Swp[ ||||| 0K/7.72G

```

Note: No program is actively using CPU.

```
Performance counter stats for 'system wide':

    2.033.097.949      cycles                    #    0,017 GHz
      55.126.405      cache-misses
    820.373.624        instructions                #    0,40  insn per cycle
    120.313,74 msec    cpu-clock                  #    4,000 CPUs utilized
      22.942          page-faults                 #    0,191 K/sec
      148.889         r01C7                       #    0,001 M/sec

    30,078463556 seconds time elapsed
```

Note: No program is actively issuing FLOPs.

Note: There are few page faults.

StressPP: Validation

Stressing CPU: What we expect to see.

We'll run *StressPP* generating a CPU workload, without specifying a CPU affinity matrix.

What we expect?

- Every core should get equally loaded.
- Few page faults.
- Increase in executed floating point instructions measured via *perf*.

Which StressPP options are used?

```
./stresspp -c 4
```

StressPP: Validation

Stressing CPU: Validation:

```
1 [|||||97.2%]
2 [|||||97.3%]
3 [|||||97.2%]
4 [|||||97.3%]
Mem[|||||3.07G/15.4G]
Swp[|||||0K/7.72G]
```

Performance counter stats for './stresspp -c 4':

356.901.771.231	cycles	#	3,404 GHz
2.217.526	cache-misses		
134.963.606.884	instructions	#	0,38 insn per cycle
104.851,34 msec	cpu-clock	#	3,495 CPUs utilized
156	page-faults	#	0,001 K/sec
268.543.578	r01C7	#	2,561 M/sec

30,002099367 seconds time elapsed

VALID: As we expected CPUs are equally loaded, there are few page faults and there's a strong increase in *FLOPs wrt idle case*.

StressPP: Validation

Stressing CPU: Validation **with affinity set on CPUs 0 and 3**

What we expect?

Same results as before but just core 0 and 3 should be loaded.

Which StressPP options are used?

```
./stresspp -cpu-affinity 1000,0001
```

StressPP: Validation

Stressing CPU: Validation **with affinity set on CPUs 0 and 3**

```
1 [|||||||||||||||||||||||||||||||||||||||||||||||||||||||||||||100.0%]
2 [|||||4.0%]
3 [||1.3%]
4 [|||||||||||||||||||||||||||||||||||||||||||||||||||||||||||||100.0%]
Mem[|||||||||||||||||||||||||||||||||||||||||3.19G/15.4G]
Swp[|||||0K/7.72G]
```

Performance counter stats for './stresspp --cpu-affinity 1000,0001':

209.396.165.980	cycles	#	3,491 GHz
1.639.193	cache-misses		
95.476.187.540	instructions	#	0,46 insn per cycle
59.983,67 msec	cpu-clock	#	1,999 CPUs utilized
155	page-faults	#	0,003 K/sec
58.341.963	r01C7	#	0,973 M/sec

30,002091388 seconds time elapsed

VALID: As we expected just CPUs **0 and 3** are stressed, there are few page faults and there's a strong increase in *FLOPs wrt idle case*.

StressPP: Validation

Stressing Virtual Memory: What we expect to see.

We'll run *StressPP* generating a Virtual Memory workload by allocating and randomly accessing 24x500Mb chunks via malloc() using two threads.

What we expect?

- Two cores should be busy performing memory operations.
- Page faults due to swap.
- High memory allocation.
- High cache misses.

Which StressPP options are used?

```
./stresspp -m 2 --vm-chunks 12 --vm-bytes 536870912
```

StressPP: Validation

Stressing Virtual Memory: Validation.



Performance counter stats for './stresspp -m 2 --vm-chunks 12 --vm-bytes 536870912':

208.969.861.415	cycles	#	3,490	GHz
5.765.460.988	cache-misses			
81.158.257.288	instructions	#	0,39	insn per cycle
59.883,06 msec	cpu-clock	#	1,965	CPUs utilized
3.146.112	page-faults	#	0,053	M/sec
0	r01C7	#	0,000	K/sec

30,471185401 seconds time elapsed

VALID: As we expected 2 CPUs are busy, *page-faults* as well as *cache-misses* are increased.

StressPP: Validation

Stressing Virtual Memory: Random Access Pattern.

In order to prove that the random access pattern is quite effective, the same experiment is performed but using a linear access pattern to read the allocated chunks

```
Performance counter stats for './stresspp -m 2 --vm-chunks 12 --vm-bytes 536870912':

198.546.072.035      cycles                #    3,473 GHz
   825.979.136      cache-misses
199.168.523.747      instructions          #    1,00  insn per cycle
   57.166,55 msec   cpu-clock                #    1,878 CPUs utilized
   3.146.110        page-faults              #    0,055 M/sec
         0          r01C7                  #    0,000 K/sec

30,442290875 seconds time elapsed
```

VALID: The *cache-misses* counter using a linear access pattern is about 7x less than the one achieved with a *random* access.



StressPP: Validation

Stressing Disk: What we expect to see.

We'll run *StressPP* continuously writing and deleting 1Gig file on hardisk using a single thread.


What we expect?

- A single core should be busy with memory operations.
- High disk usage.

Which StressPP options are used?

```
./stresspp -d 1
```

Stressing Disk: Validation.



```
104.528.446.847      cycles                #      3,488 GHz
    1.419.306.550      cache-misses
143.560.850.735      instructions          #      1,37  insn per cycle
    29.966,59 msec     cpu-clock                #      0,996 CPUs utilized
        411           page-faults             #      0,014 K/sec
            0          r01C7                   #      0,000 K/sec

    30,087785071 seconds time elapsed
```

Total DISK READ :		0.00 B/s	Total DISK WRITE :		1731.79 M/s		
Actual DISK READ:		0.00 B/s	Actual DISK WRITE:		0.00 B/s		
TID	PRI0	USER	DISK READ	DISK WRITE	SWAPIN	IO>	COMMAND
12464	be/4	root	0.00 B/s	732.48 M/s	0.00 %	0.00 %	./stresspp -d 1
1	be/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	init
2	be/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[kthreadd]

VALID: As we expected there's a single core busy in writing while the disk is written at *700Mb/s*



StressPP: Validation

Stressing I/O: What we expect to see.

We'll run *StressPP* continuously asking to sync() using a single thread.

What we expect?

- High I/O requests.
- Low CPU usage since there are no operations in charge of the processor.

Which StressPP options are used?

```
./stresspp -i 1
```

StressPP: Validation

Stressing I/O: What we expect to see.

```
1 [||||| 14.2%]
2 [||||| 12.1%]
3 [||||| 16.3%]
4 [||||| 16.5%]
Mem[||||| 3.49G/15.4G]
Swp[||||| 999M/7.72G]
```

Performance counter stats for './stresspp -i 1':

19.283.335.563	cycles	#	2,956 GHz
16.026.653	cache-misses		
22.322.442.337	instructions	#	1,16 insn per cycle
6.524,10 msec	cpu-clock	#	0,217 CPUs utilized
154	page-faults	#	0,024 K/sec
0	r01C7	#	0,000 K/sec

30,002050080 seconds time elapsed

Total DISK READ :	0.00 B/s	Total DISK WRITE :	15.66 K/s
Actual DISK READ:	0.00 B/s	Actual DISK WRITE:	31.32 K/s

TID	PRIO	USER	DISK READ	DISK WRITE	SWAPIN	IO>	COMMAND
16032	be/4	root	0.00 B/s	0.00 B/s	0.00 %	65.10 %	./stresspp -i 1
312	be/3	root	0.00 B/s	3.91 K/s	0.00 %	0.46 %	[ibd2/dm-0-8]

VALID: As we expected there's no busy core and *stresspp* is actually the most I/O hungry program.



Thank you.

<https://github.com/Guglio95/StressPP>